


From Paper to Practice: Leveraging LST-Bench to Evaluate Lake-Centric Data Platforms

Jesús Camacho-Rodríguez
Gray Systems Lab (GSL)
Microsoft Azure Data

DBTest '24 Keynote
June 9th, 2024



① Data Lake-Centric Architectures and LSTs

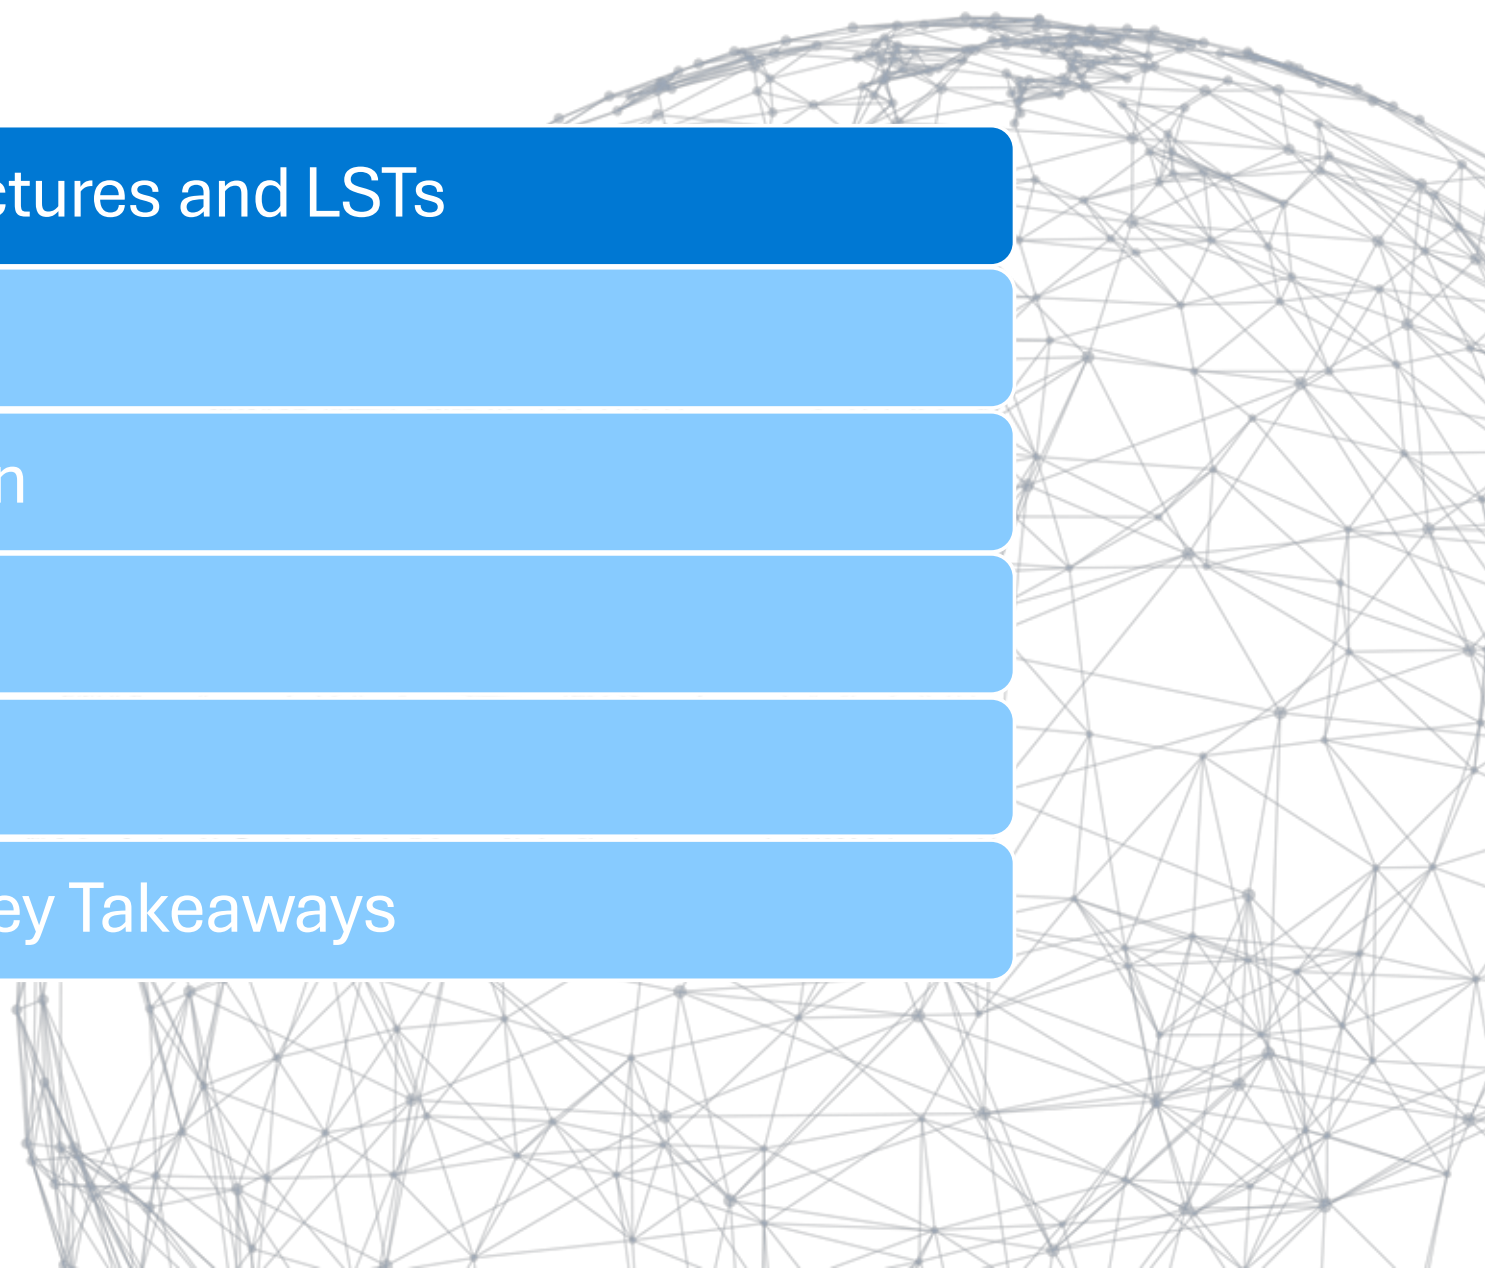
② Motivation

③ Tool and Benchmark Design

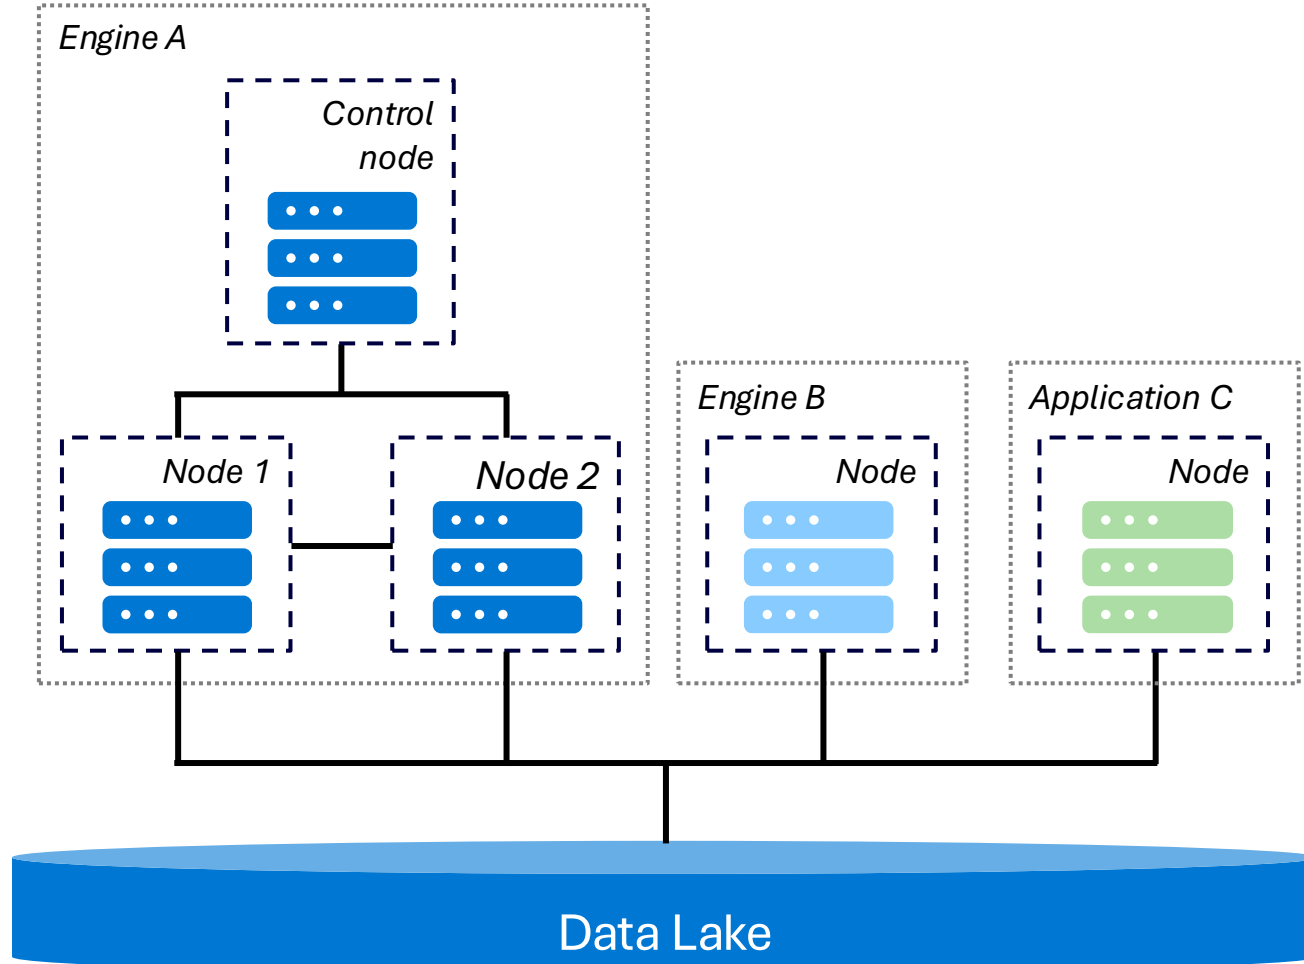
④ Implementation

⑤ Results

⑥ Status, Road Ahead, and Key Takeaways



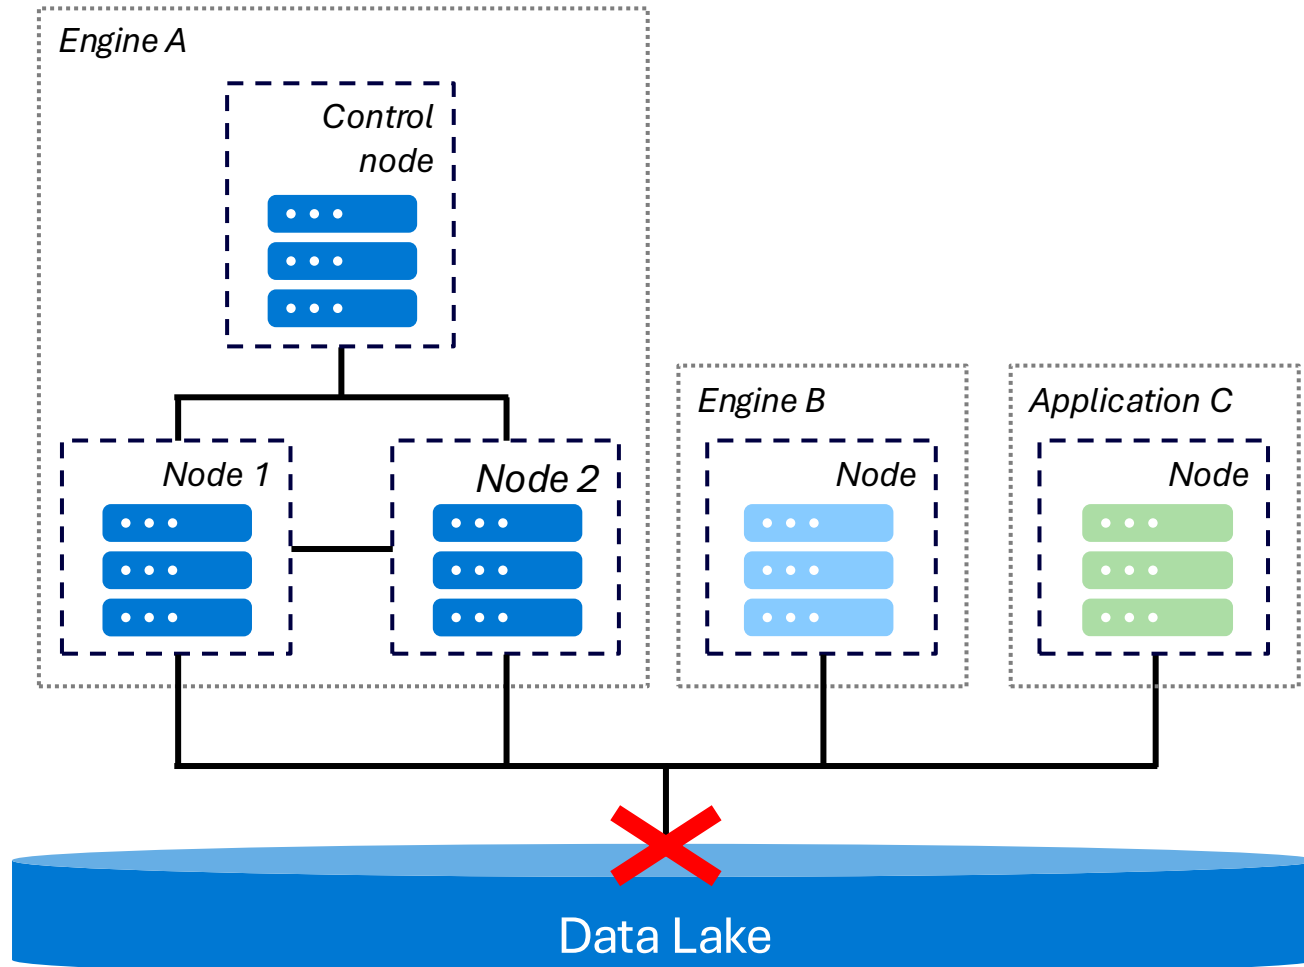
Data Lake-Centric Architecture



Benefits

- **Flexible Scalability:**
Scale storage and compute independently for better efficiency and cost savings
- **Streamlined Workflows:**
Eliminate data silos for simpler data movement across systems
- **Reduced Lock-In:**
Allow any engine to directly access data storage, offering flexibility to choose the best solution for each application

Data Lake-Centric Architecture



Challenges

- Ensuring consistency and isolation of complex read and write transactions
- Data lakes excel in scalability and durability but lack necessary concurrency and recovery capabilities

Log Structured Tables (LSTs) – Overview



- Popular OSS projects for updatable tables: **Delta Lake**, **Apache Hudi**, **Apache Iceberg**, and **Apache Paimon (Incubating)**
 - **Goal:** Provide additional functionality (transactions, indexes, time travel, cloning) on top of immutable files (Parquet) stored in the data lake



LSTs – Data and Protocol



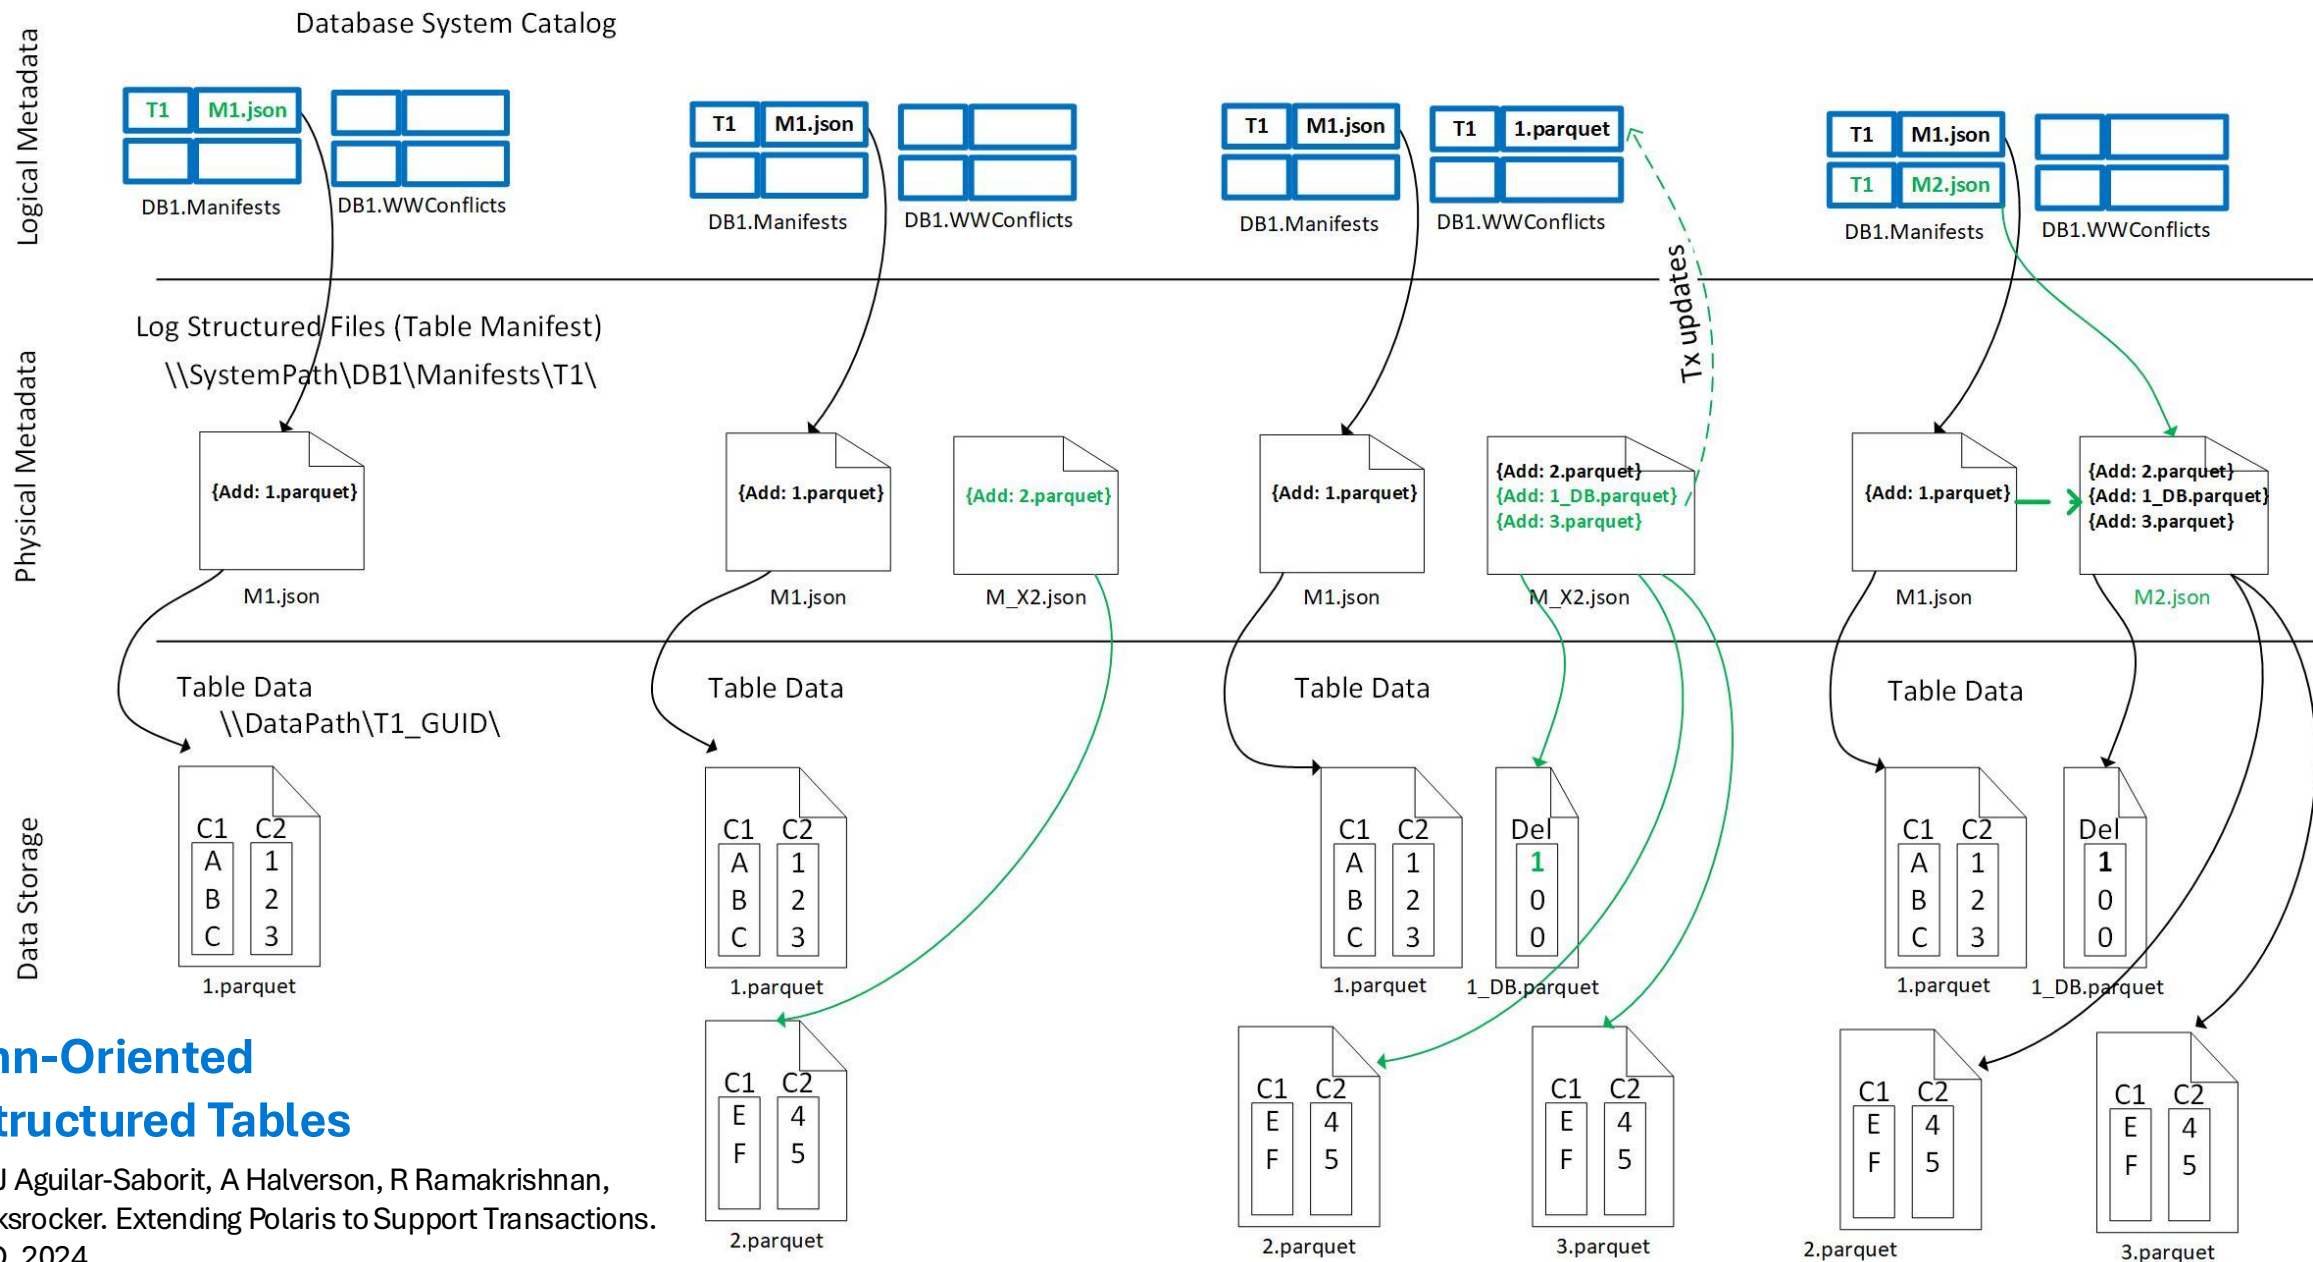
- Many types of data are required to represent a LST
 - Logical metadata (e.g., table definitions)
 - Physical metadata (how to find data for a table)
 - Data files (Parquet type, column oriented, store user data)
 - Delete bitmaps (describes which rows are deleted)
 - Clone info (reference counting for files)
- Where this data is stored
 - Files in cloud storage
- Approach to updates — NOT in-place in page-files!
 - Copy-on-Write (CoW) — Affected data files copied over with changes reflected
 - Merge-on-Read (MoR) — For affected data files, ONLY changes recorded; reads must merge changes

(1) X1. Loaded 3 rows
and committed

(2) X2 bulk inserts 2 rows

(3) X2 deletes one row from 1.parquet
X2 inserts 2 more rows

(4) X2 commits



Column-Oriented Log-Structured Tables

[AHR+24] J Aguilar-Saborit, A Halverson, R Ramakrishnan,
and K Bocksrocker. Extending Polaris to Support Transactions.
In SIGMOD, 2024.



① Data Lake-Centric Architectures and LSTs

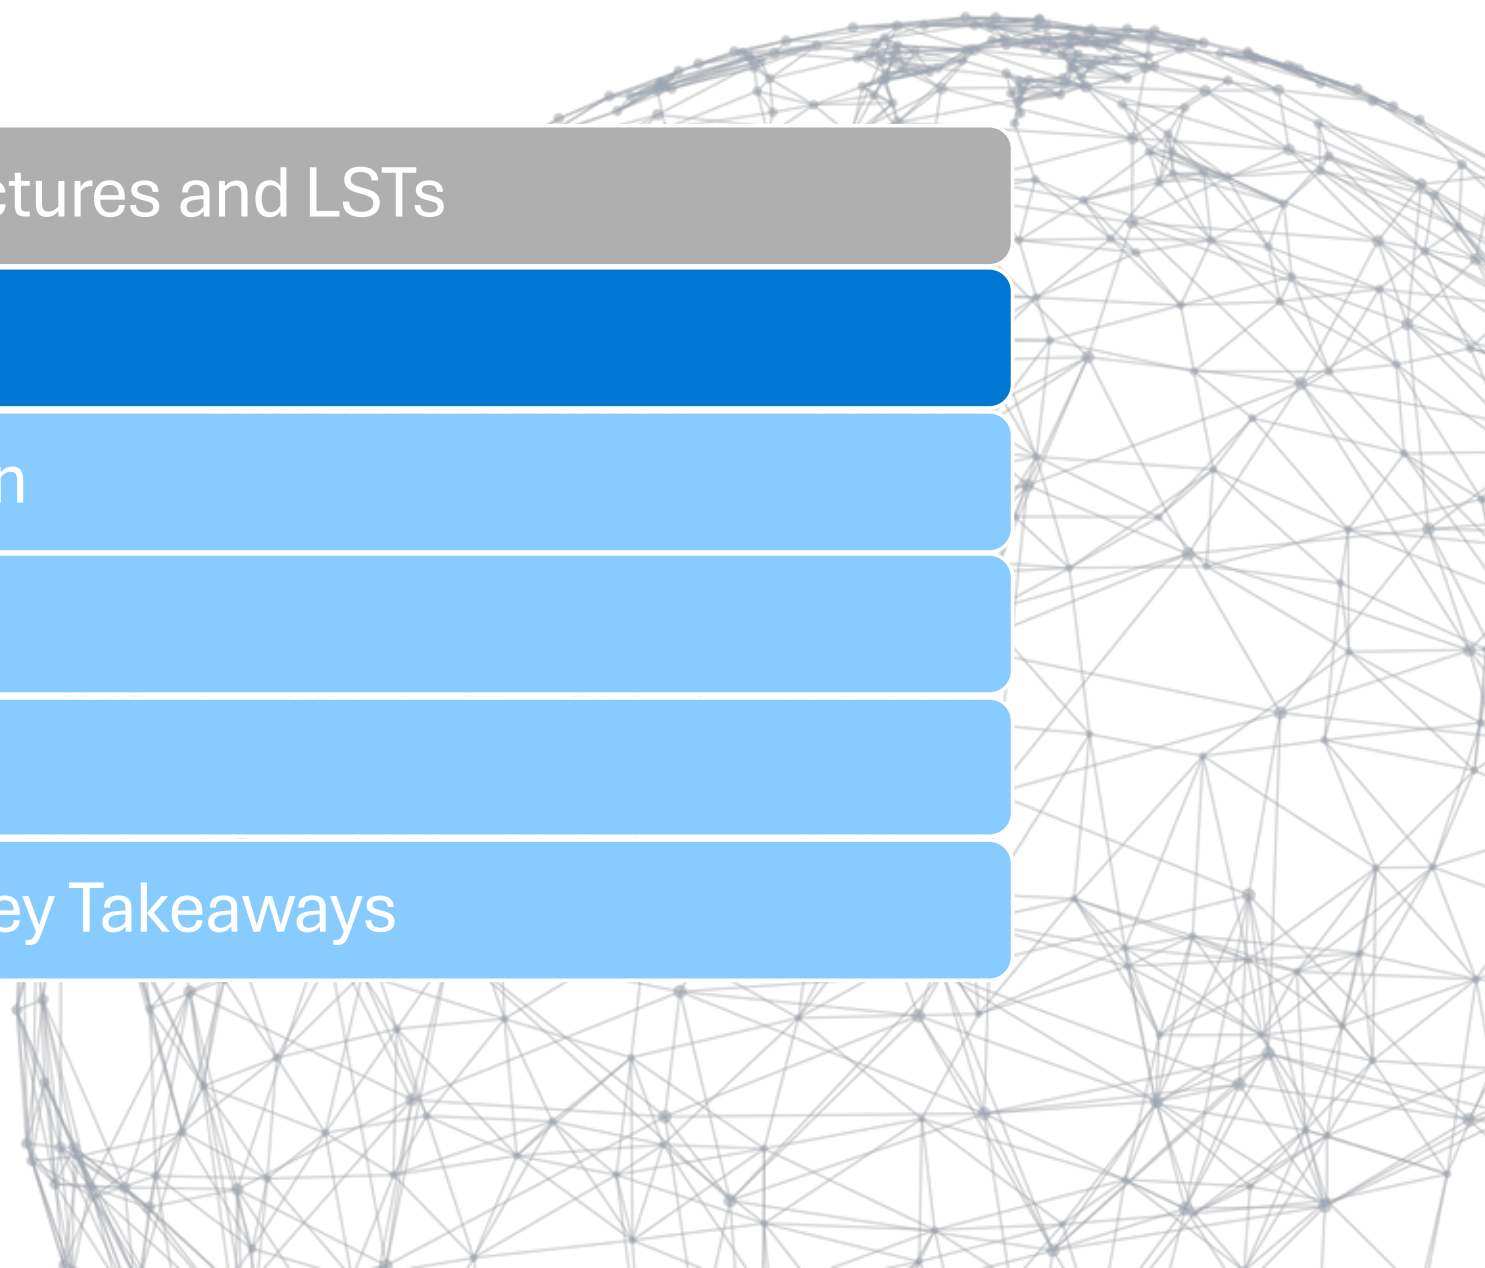
② Motivation

③ Tool and Benchmark Design

④ Implementation

⑤ Results

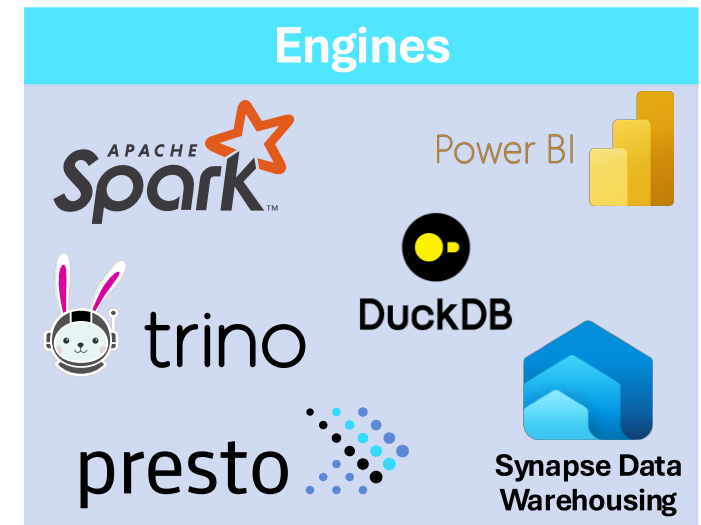

⑥ Status, Road Ahead, and Key Takeaways



LSTs, Engines, and Platforms



- While LSTs share the same goal, their architecture and implementation vary
 - CoW and MoR support, metadata caching, distributed planning, storage optimizations
- Performance depends on numerous factors



- Limited work on a **comprehensive framework for LSTs evaluation**

Existing Approaches to Evaluate LSTs



- **Experimental Evaluation:** LH-Bench [JKP+23], Brooklyn Data, DataBeans
 - Ad-hoc approaches
 - Typically rely on **TPC-DS**, standard OLAP benchmark
 - Subset of standard queries or handcrafted queries

Limitation 1: Unable to extend to new engines, datasets, and scenarios beyond traditional OLAP tasks, despite the continuous expansion of data-lake centric architectures with new use cases

TPC-DS Overview

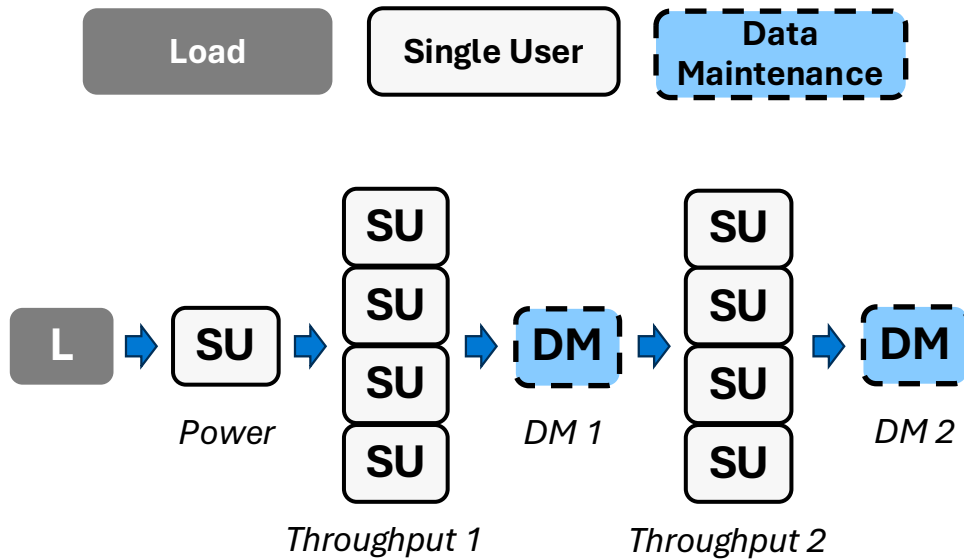


Workload:

Up to 99 analytical queries ?

Performance Metric:

Total execution latency ?



Normalized *query throughput per hour*

Product of total # of queries executed
and scale factor


$$Q_{phDS@SF} = \left[\frac{SF * Q}{\sqrt[4]{T_{PT} * T_{TT} * T_{DM} * T_{LD}}} \right]$$

Geometric mean of elapsed time for
load, power, throughput, and data
maintenance phases


TPC-DS Workload



Question: Regarding trickle updates, is TPC-DS a representative workload?



Limitation 2: Failure to expose important characteristics of LSTs that are crucial in real customer scenarios




Load

Single User

Data
Maintenance

TPC-DS Metrics



Limitation 3: Lack of metrics to expose important aspects when comparing different LST implementations

$$QphDS@SF = \left\lfloor \frac{SF * Q}{\sqrt[4]{T_{PT} * T_{TT} * T_{DM} * T_{LD}}} \right\rfloor$$

LST	Throughput-QphDS	Inter-test Degradation
Delta	511K	2.7 -> 5.2 hrs (92%)
Hudi-CoW*	262K	6.2 -> 6.5 hrs (5%)
Hudi-MoR*	112K	23 -> 24 hrs (6%)
Iceberg-CoW*	549K	2.7 -> 4 hrs (45%)
Iceberg-MoR*	493K	2.9 -> 5 hrs (73%)

* Copy-on-Write mode

* Merge-on-Read mode

SF=1000

Spark (3.3.1), 16 workers, 8 cores, 64 GB

Delta (2.2.0), Iceberg (1.1.0), Hudi (0.12.2)

Load

Single User

Data
Maintenance



① Data Lake-Centric Architectures and LSTs

② Motivation

③ Tool and Benchmark Design

④ Implementation

⑤ Results

⑥ Status, Road Ahead, and Key Takeaways

LST-Bench



- A benchmark tool specifically tailored to evaluate LSTs

Limitation 1: Unable to extend to new engines, datasets, and scenarios beyond traditional OLAP tasks, despite the continuous expansion of data-lake centric architectures with new use cases



Flexible Workload Representation

Can represent existing benchmarks but easily extends to new scenarios

Limitation 2: Failure to expose important characteristics of LSTs that are crucial in real customer scenarios



Enhanced Workloads

Introduces *diverse workload patterns*

Limitation 3: Lack of metrics to expose important aspects when comparing different LST implementations



Relevant Metrics

Includes *metrics extensions* relevant for LSTs

Workload Representation

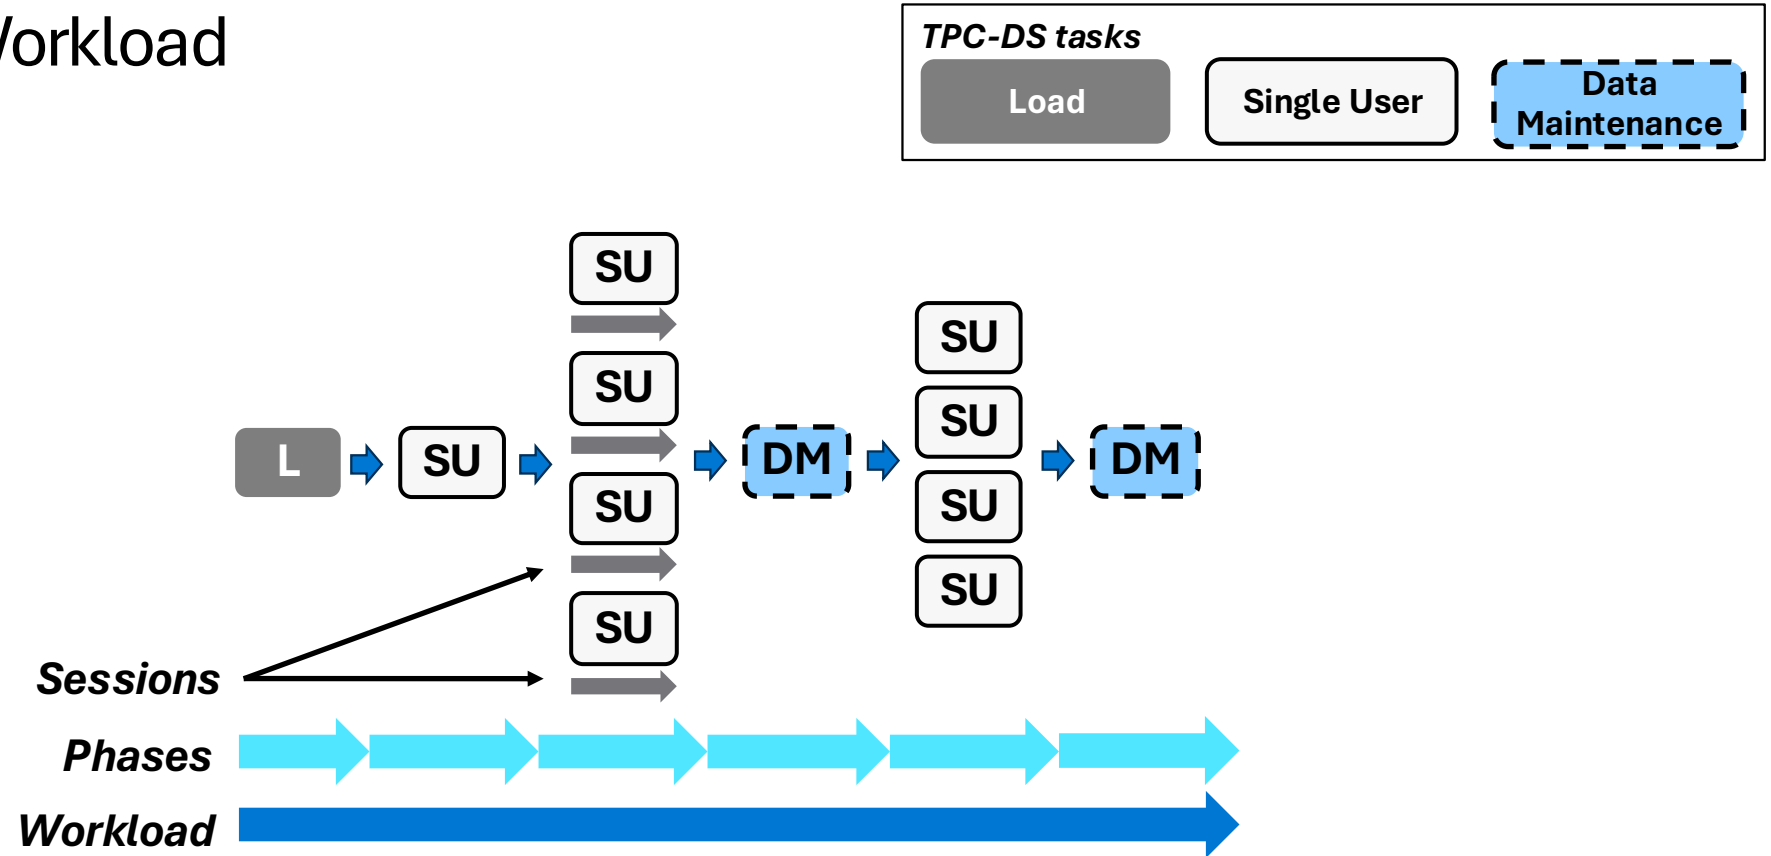


- **Adaptable to both standard and custom workloads prevalent in practice**
 - Driven by diverse and extensive internal user feedback
 - Facilitates the mapping of existing workloads, such as TPC-DS
 - Aligns with the concept of sessions in JDBC
 - Ensures ease of reusability

Extending Workloads Beyond TPC-DS



- **W0:** TPC-DS Workload



LST-Bench Workload Patterns

TPC-DS tasks

Load

Single User

Data
Maintenance

New tasks

Optimize

Gain insights into LST aspects overlooked by TPC-DS


- W1: How are frequent data modifications handled over a long period of time?



- W2: How are multiple data modifications of varying sizes handled within a regularly optimized table?



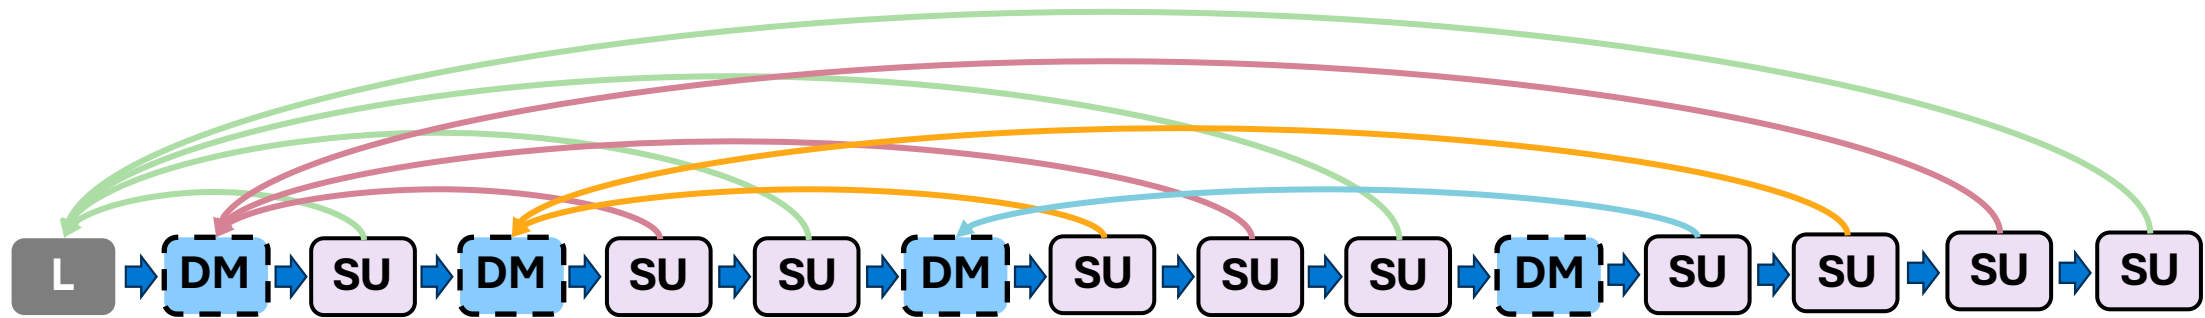
- W3: How are simultaneous reading and writing sessions handled, potentially across multiple compute clusters?



LST-Bench Workload Patterns

Gain insights into LST aspects overlooked by TPC-DS

- **W4:** How is data querying across different points in time handled?



- **W5:** How does updating or deleting the same data volume in different batch sizes influence read query performance?



*n tasks dynamically generated
by parameterized custom task*

TPC-DS tasks

Load

Single User

Data
Maintenance

New tasks

Optimize

Single User
(Time Travel)

Parameterized
Custom

LST-Bench Metrics



Traditional Metrics

- **Performance:** *Latency, Throughput*
- **Cloud Storage Efficiency:** *Capacity Utilization, API Call Count, Total IO*
- **Compute Engine Efficiency:** *CPU Utilization, Memory Utilization, Disk Utilization*


Other Metrics

- **Stability** - System's ability to sustain consistent performance and efficiency with minimal degradation: *Degradation Rate*

$$S_{DR} = \frac{1}{n} \sum_{i=1}^n \frac{M_i - M_{i-1}}{M_{i-1}}$$

where

- M_i is metric value of the i^{th} iteration of a workload phase,
- n is the number of iteration of the phase, and
- S_{DR} is the degradation rate.





① Data Lake-Centric Architectures and LSTs

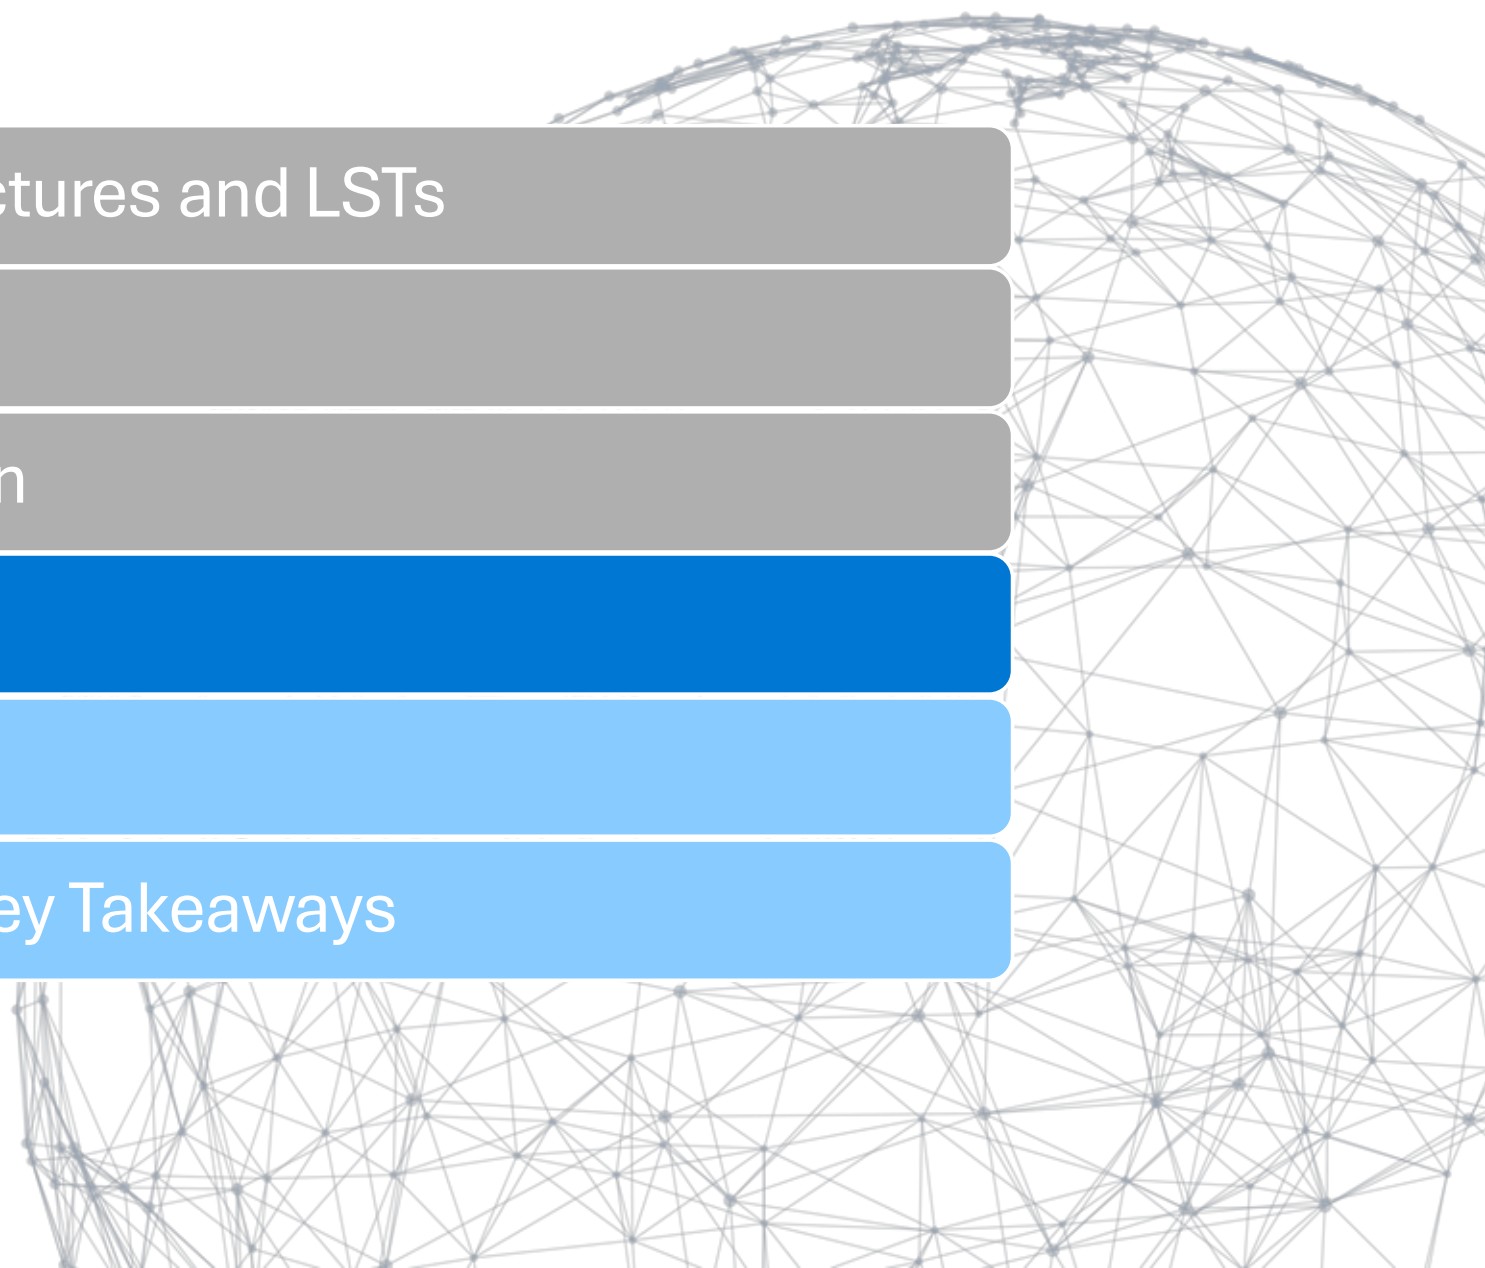
② Motivation

③ Tool and Benchmark Design

④ Implementation

⑤ Results

⑥ Status, Road Ahead, and Key Takeaways

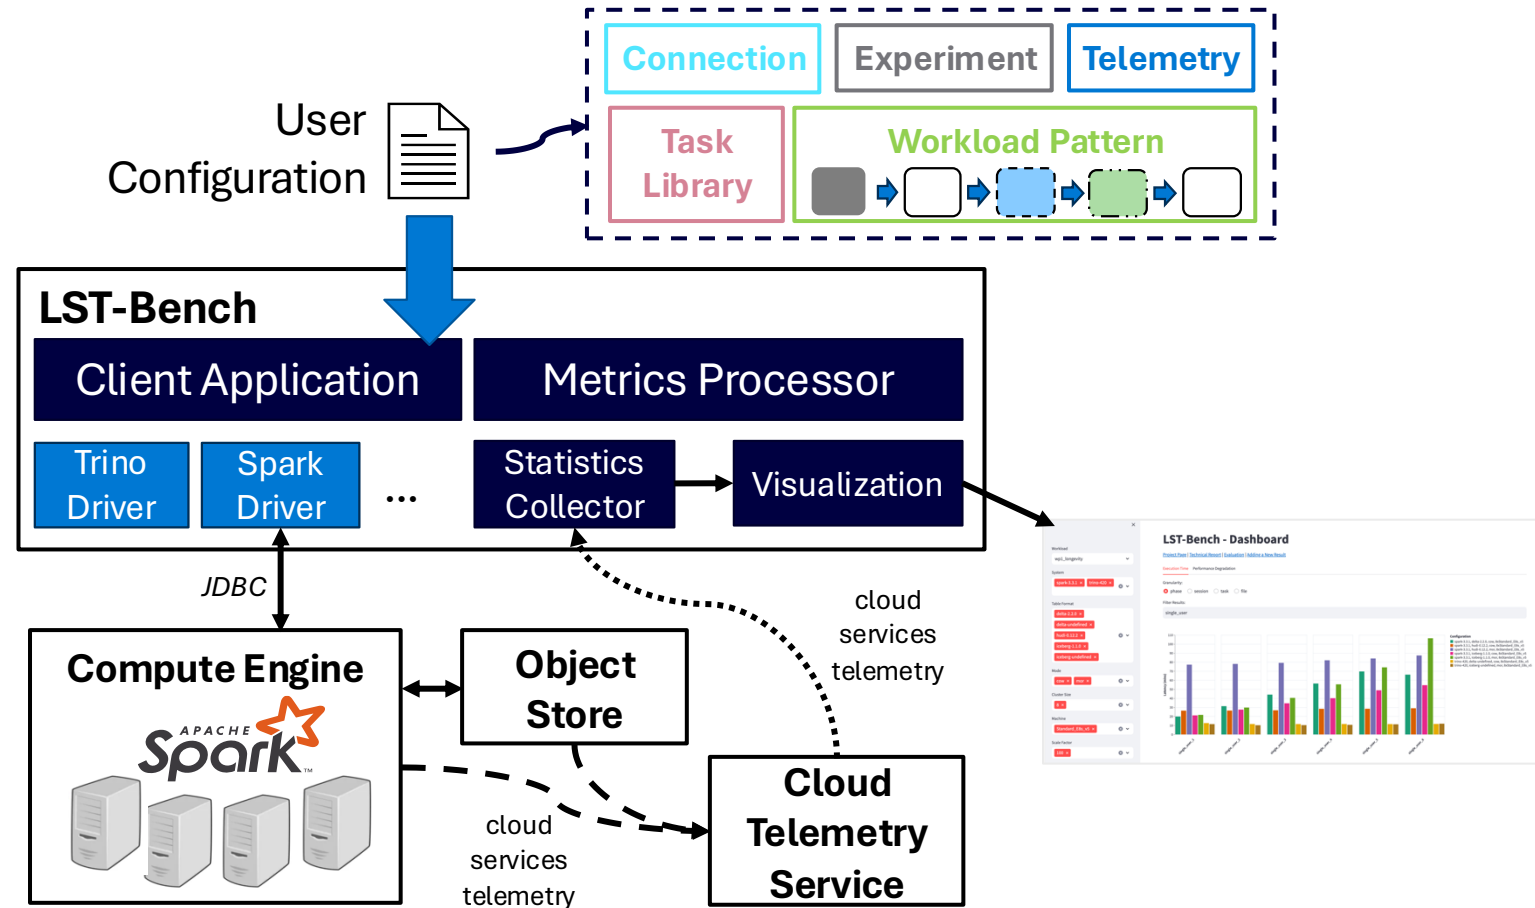


LST-Bench Implementation



- Java Client Application
 - Customizable / extensible via config files
 - Connects to engine via JDBC or Spark session client
- Python Metrics Processor
 - Visualization via notebook or Streamlit web app
- Open-source available under Apache License 2.0:

<https://github.com/microsoft/lst-bench>



LST-Bench Configuration, Libraries, and Workloads



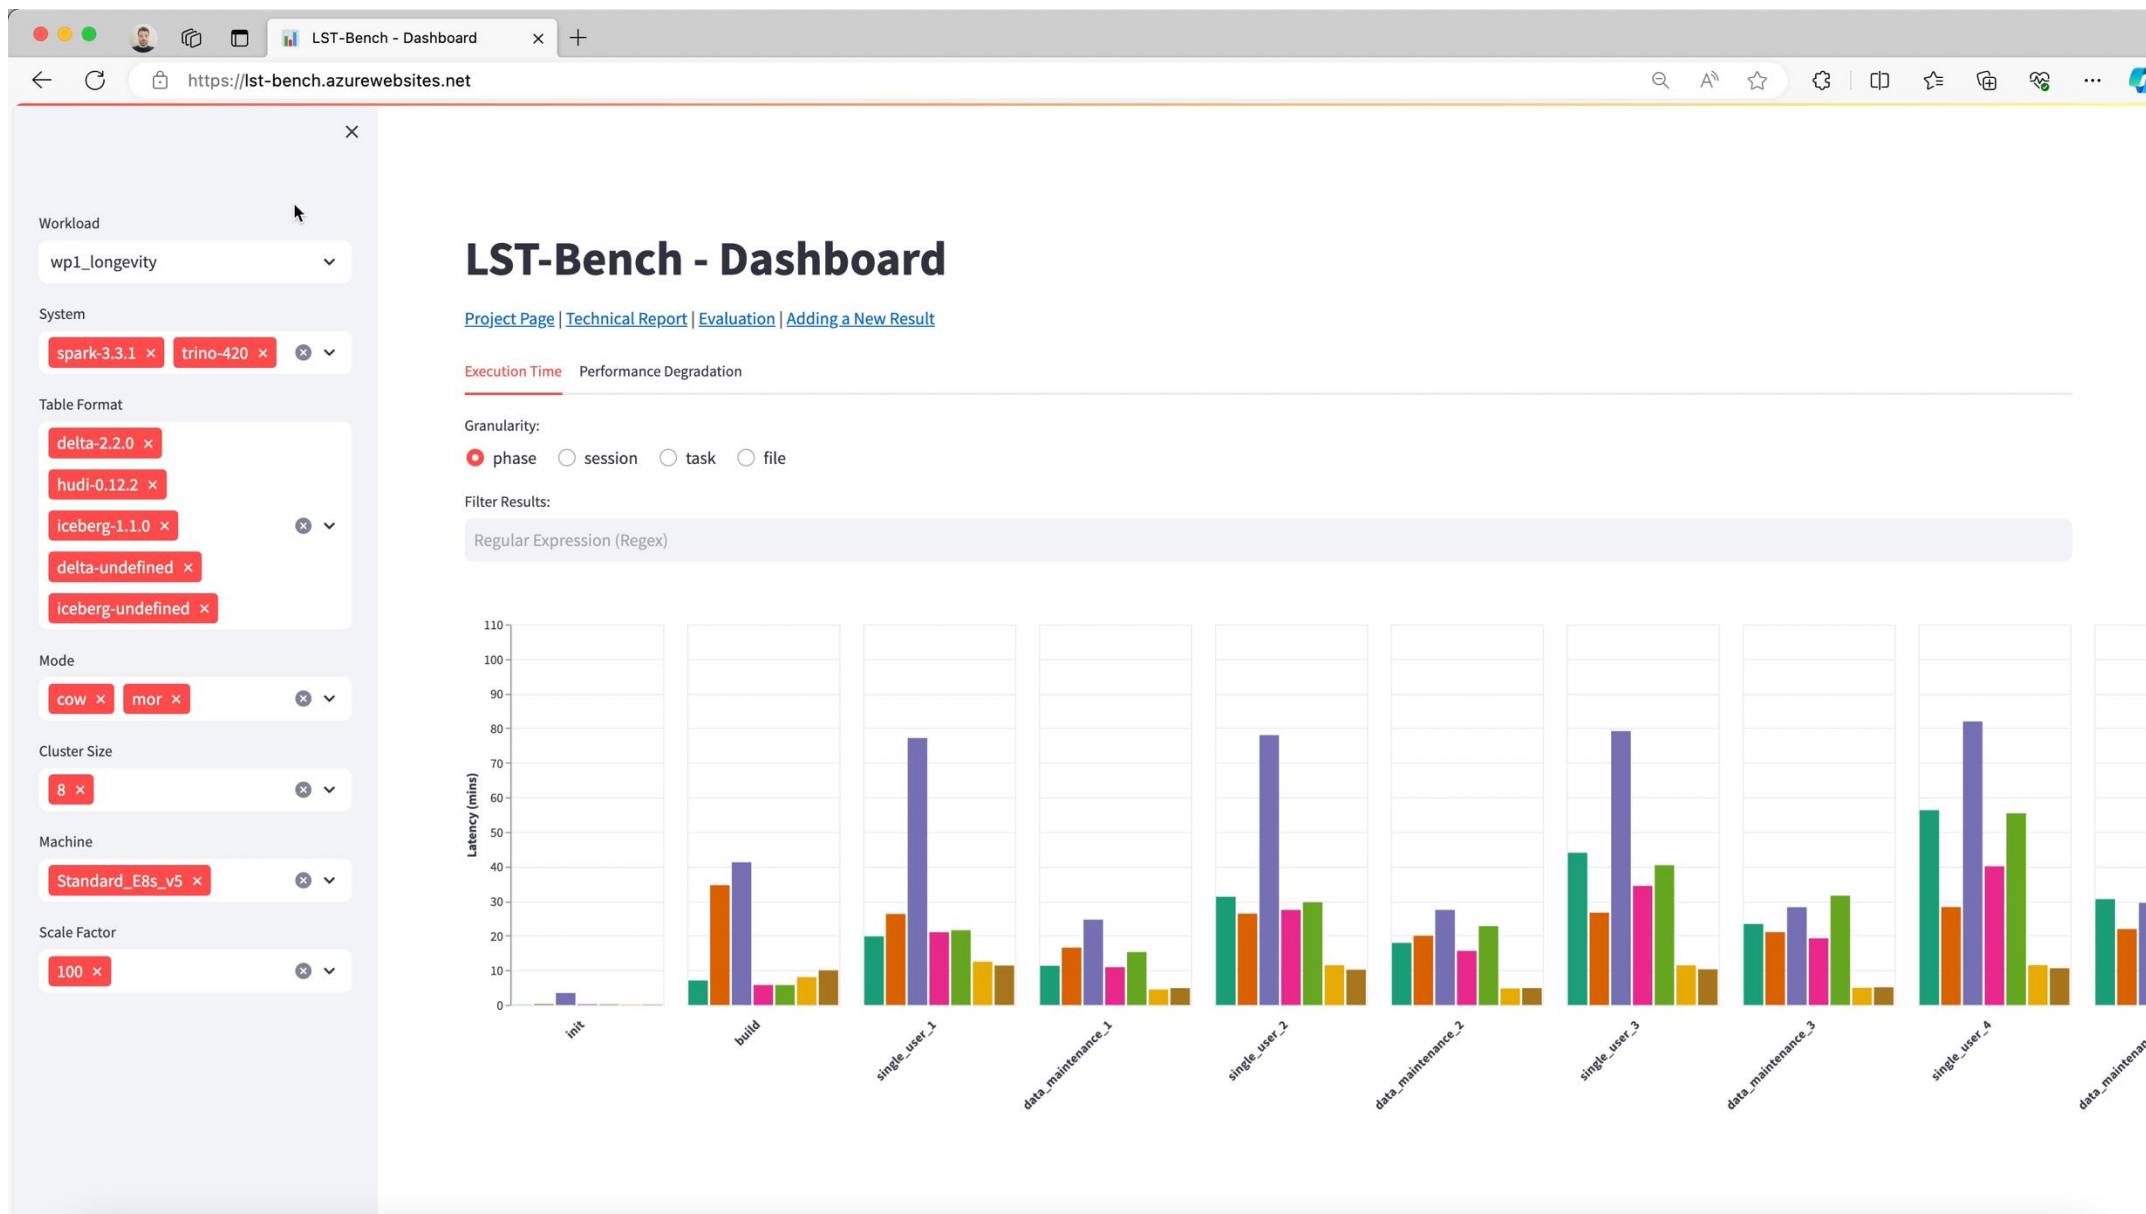
- **Configuration-driven approach**
 - YAML/JSON file support
 - Input validation against JSON Schema
- **Workload options**
 - **Self-contained:** All tasks, sessions, and phases defined within the workload file
 - **Library-based:** Define tasks, sessions, and phases in a library and reference these entities from the workload definition

```
- id: throughput_simple_phase
sessions:
- tasks:
  - template_id: single_user_simple
    permute_order: true
  target_endpoint: 0
- tasks:
  - templ
    permu
  target_
- tasks:
  - templ
    permu
  target_
- tasks:
  - templ
    permu
  target_

id: my_first_workload
phases:
- id: warm_up
  sessions:
  - tasks:
    - template_id: single_user_simple
      target_endpoint: 0
  - tasks:
    - template_id: single_user_simple
      target_endpoint: 1
- id: throughput_simple
  template_id: throughput_simple_phase
```

workload.yaml

LST-Bench Web UI





① Data Lake-Centric Architectures and LSTs

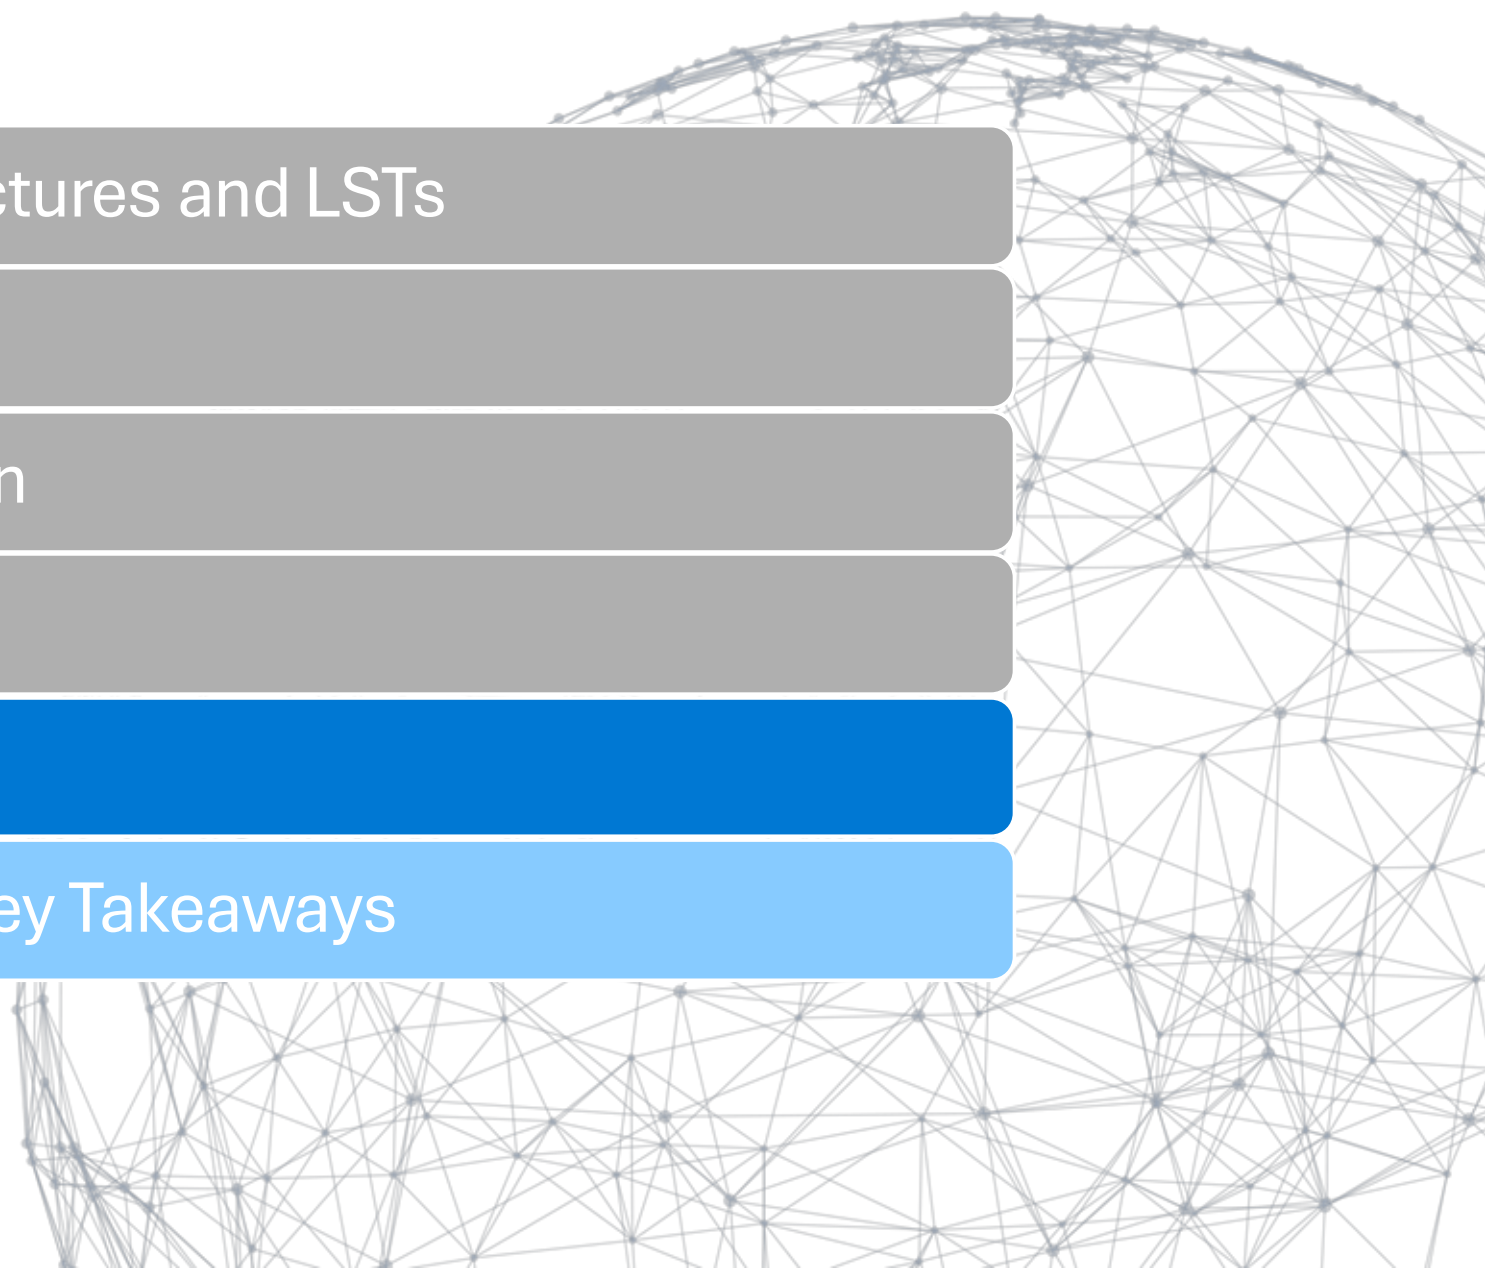
② Motivation

③ Tool and Benchmark Design

④ Implementation

⑤ Results

⑥ Status, Road Ahead, and Key Takeaways



Evaluation



- Azure VMSS cluster with 1 head and 16 worker nodes
 - Each node with 8 virtual cores and 64GB RAM
- Data stored in Azure Data Lake Storage Gen2 (ADLS)
 - TPC-DS SF1000
- Azure Monitor to collect telemetry and Logs Analytics to execute queries against it
- No special tuning for any of the engines, LSTs we evaluated:
 - [Apache Spark 3.3.1](#): Delta Lake v2.2.0, Apache Hudi v0.12.2, Apache Iceberg v1.1.0
 - [Trino 420](#): Delta Lake, Apache Iceberg
- [Important remarks](#)
 - Results subject to change and improvements due to further tuning and future developments
 - Insights drawn for these engines may not apply to the LST on different engines

Evaluation – *W1*




Performance (Spark): Significant slowdown observed across iterations (up to 6.8x) for Iceberg-MoR. Nearly all formats show a decrease in performance.

Evaluation – *W1*




Performance (Trino compared to Spark): Nearly double the speed for both Delta and Iceberg tables after load. Higher stability.

Evaluation – *W1*





Storage Layer Calls: Increased API calls due to a higher number of (small) files that need to be read.

Evaluation – *W1*



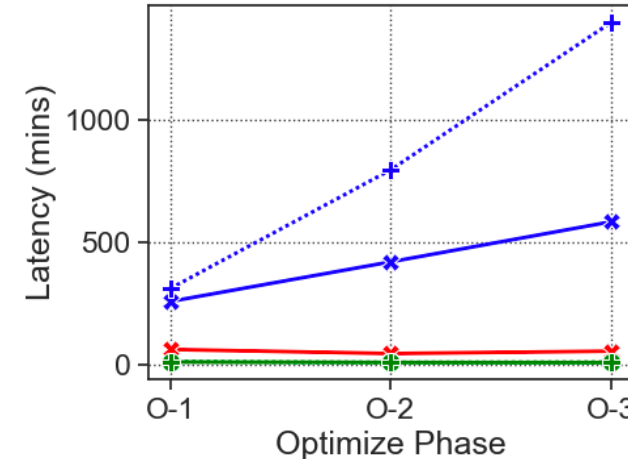
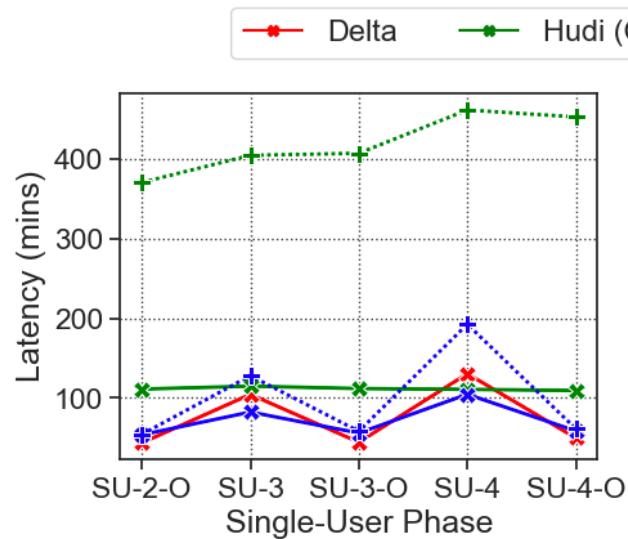

I/O Volume: Comparatively high I/O volume for Hudi due to its default configuration, resulting in a higher number of small files.

Evaluation – *W2*





1. Table maintenance has a big impact on Delta and Iceberg performance stability (*zig-zag pattern*).
2. Hudi maintains stable performance without user-triggered maintenance by doing work upfront.

Evaluation – *W2*





Iceberg's default file grouping for compaction significantly increases compaction time (potentially minimizes read query disruptions) → Tuning LSTs involves trade-offs based on user goals

Evaluation – *W2 vs W3*



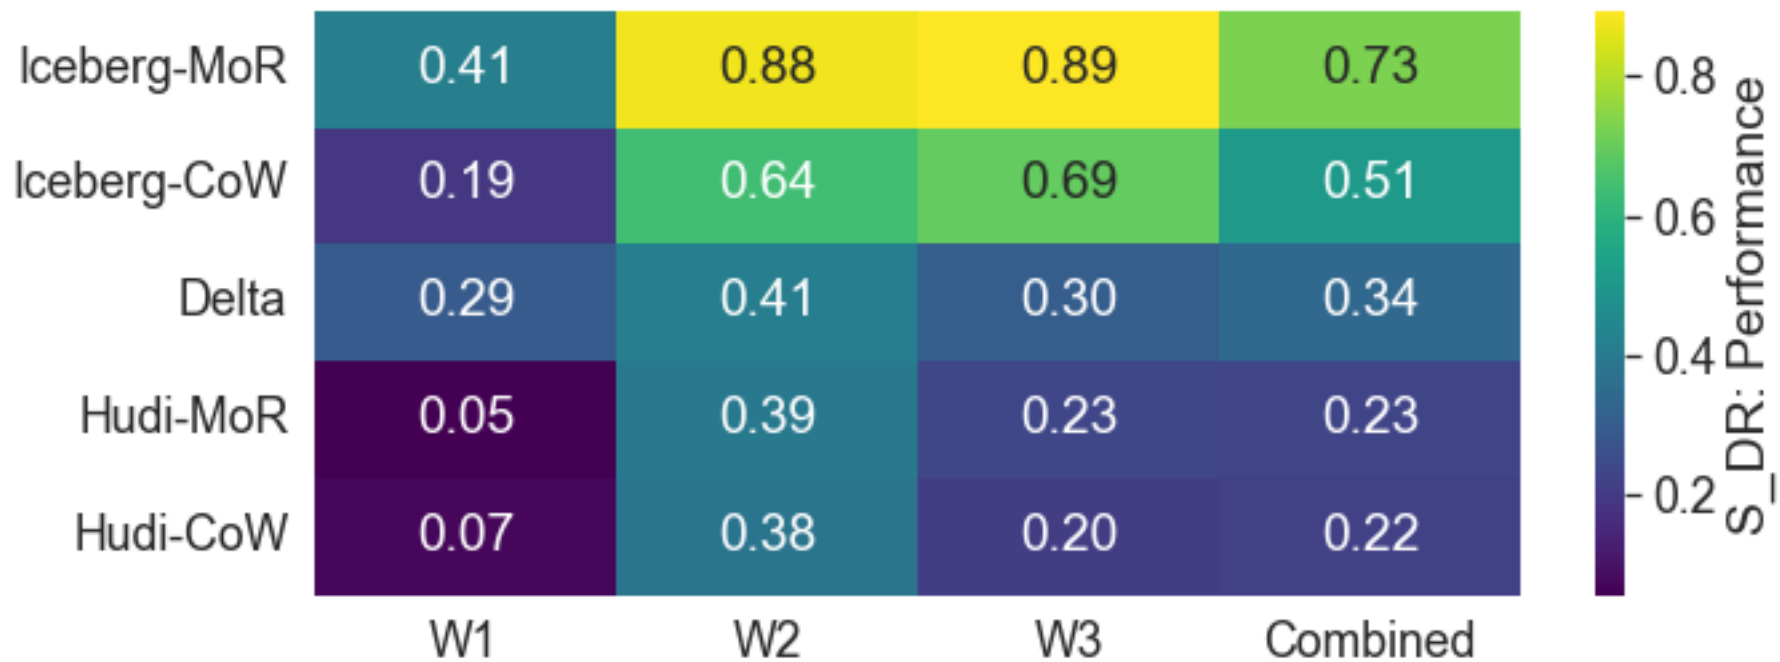
Spark: Concurrent session execution does not lead to significant performance improvements due to resource contention.

Evaluation – *W2 vs W3*




Trino: Efficient utilization of cluster resources results in significant end-to-end experiment runtime gains, despite minor slowdowns in individual statements.

Stability Evaluation



Stability Evaluation




Hudi shows highest stability

Stability Evaluation



Iceberg shows lowest stability



Hudi shows highest stability



① Data Lake-Centric Architectures and LSTs

② Motivation

③ Tool and Benchmark Design

④ Implementation

⑤ Results

⑥ Status, Road Ahead, and Key Takeaways

Status and Road Ahead




- **LST-Bench @ Microsoft**

- Integrated into Microsoft Fabric Warehouse's testing workflow
- Foundational tool for various ongoing initiatives:
 - Automatic tuning and maintenance policies for LSTs
 - Performance evaluation of LSTs converted using [Apache XTable \(Incubating\)](#)

- **Open-source LST-Bench**

- Support for other engines, platforms, cloud providers (Apache Flink, Snowflake, AWS)
- New scenarios: Data cleaning, CDC with transactional consistency guarantees
- Integration with OpenTelemetry

- **Others? Contributions welcome!**



Key Takeaways



- **Evolving Benchmarks:**

- Traditional OLAP benchmarks like TPC-DS are not representative of modern analytic data lake workloads, e.g., lack of trickle updates

- **Flexible and Extendable Tools:**

- **Modular, flexible benchmarking tools** are essential for evaluating new engines, datasets, and scenarios in the ever-expanding landscape of data lake architectures

- **Comprehensive Metrics and Observability:**

- One representative metric can simplify decision-making, but **enhanced metrics and 360-degree observability** are crucial for understanding system characteristics
- Achieving this level of observability is **challenging**, especially across multiple engines and cloud environments

Acknowledgements



**Ashvin
Agrawal**



**Anja
Gruenheid**



**Jose
Medrano**



**Emma Rose
Wirshing**



**Ashit
Gosalia**



**Cristian
Petculescu**



**Josep
Aguilar-Saborit**



**Avriia
Floratou**




**Carlo
Curino**



**Raghu
Ramakrishnan**



LST-Bench paper @
ACM SIGMOD 2024



Open-source available
(Apache License 2.0)

Thank you! Questions?